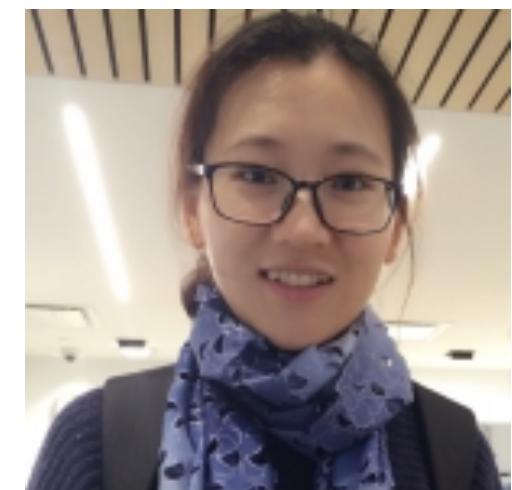


Lecture 2 (cont'd) & Lecture 3: Advanced SQL – Part I

Announcements!

1. You should be Jupyter notebook Ninjas!
2. Welcome Ting!
 - New TA-Office hours on website (room to be announced)
3. Project groups finalized!
 - If you do not have a group talk with us ASAP!
4. Problem Set #1 released



Lecture 2 (cont'd) & Lecture 3: Advanced SQL – Part I

Today's Lecture

1. Recap from Lecture 2 & Multi-table queries
 - ACTIVITY: Multi-table queries
2. Set operators & nested queries
 - ACTIVITY: Set operator subtleties

Lecture 2 (cont'd): Introduction to SQL

3. Multi-table queries

What you will learn about in this section

1. Primary keys and Foreign keys recap
2. Joins: SQL semantics
3. ACTIVITY: Multi-table queries

Keys and Foreign Keys

Company

CName	StockPrice	Country
GizmoWorks	25	USA
Canon	65	Japan
Hitachi	15	Japan

What is a foreign key vs. a key here?

Product

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

A **key** is a **minimal subset of attributes** that acts as a unique identifier for tuples in a relation

If two tuples agree on the values of the key, then they must be the **same** tuple!

Keys and Foreign Keys

Company

CName	StockPrice	Country
GizmoWorks	25	USA
Canon	65	Japan
Hitachi	15	Japan

What is a foreign key vs. a key here?

A **foreign key** is an attribute (or collection of attributes) in one table that uniquely identifies a row of another table.

Product

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

The foreign key is **defined** in a second table, but it refers to the primary key in the first table.

Declaring Foreign Keys

```
Company(CName: string, StockPrice: float, Country: string)
Product(PName: string, Price: float, Category: string, Manufacturer: string)

CREATE TABLE Product(
    pname          VARCHAR(100),
    price          FLOAT,
    category       VARCHAR(100),
    manufacturer   VARCHAR(100),
    PRIMARY KEY (pname, manufacturer),
    FOREIGN KEY (manufacturer) REFERENCES Company(cname)
)
```

Declaring Foreign Keys

```
CREATE TABLE Company(  
    cname      VARCHAR(100),  
    stockprice FLOAT,  
    country    VARCHAR(100),  
    PRIMARY KEY (cname),  
    FOREIGN KEY (cname) REFERENCES Product(pname, manufacturer)  
)
```

```
CREATE TABLE Product(  
    pname      VARCHAR(100),  
    price      FLOAT,  
    category   VARCHAR(100),  
    manufacturer VARCHAR(100),  
    PRIMARY KEY (pname, manufacturer)  
)
```

Can we do this? What would be
the problem?

Declaring Foreign Keys

```
CREATE TABLE Company(  
    cname      VARCHAR(100),  
    stockprice FLOAT,  
    country    VARCHAR(100),  
    PRIMARY KEY (cname),  
    FOREIGN KEY (cname) REFERENCES Product(pname, manufacturer)  
)
```

```
CREATE TABLE Product(  
    pname      VARCHAR(100),  
    price      FLOAT,  
    category   VARCHAR(100),  
    manufacturer VARCHAR(100),  
    PRIMARY KEY (pname, manufacturer)  
)
```

Can we do this? What would be
the problem?

*We can have products without a registered
company! Bad design! We'll see more next week.*

Declaring Foreign Keys

```
CREATE TABLE Company(
    cname          VARCHAR(100),
    stockprice     FLOAT,
    country        VARCHAR(100),
    PRIMARY KEY (cname),
    FOREIGN KEY (cname) REFERENCES Product(pname, manufacturer)
)
```

```
CREATE TABLE Product(
    pname          VARCHAR(100),
    price          FLOAT,
    category       VARCHAR(100),
    manufacturer   VARCHAR(100),
    PRIMARY KEY (pname, manufacturer)
)
```

If the **primary key** is a set of columns (a **composite key**), then the **foreign key** also must be a set of columns that corresponds to the **composite key**.

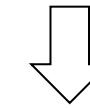
Joins

Product

PName	Price	Category	Manuf
Gizmo	\$19	Gadgets	GWorks
Powergizmo	\$29	Gadgets	GWorks
SingleTouch	\$149	Photography	Canon
MultiTouch	\$203	Household	Hitachi

Company

Cname	Stock	Country
GWorks	25	USA
Canon	65	Japan
Hitachi	15	Japan

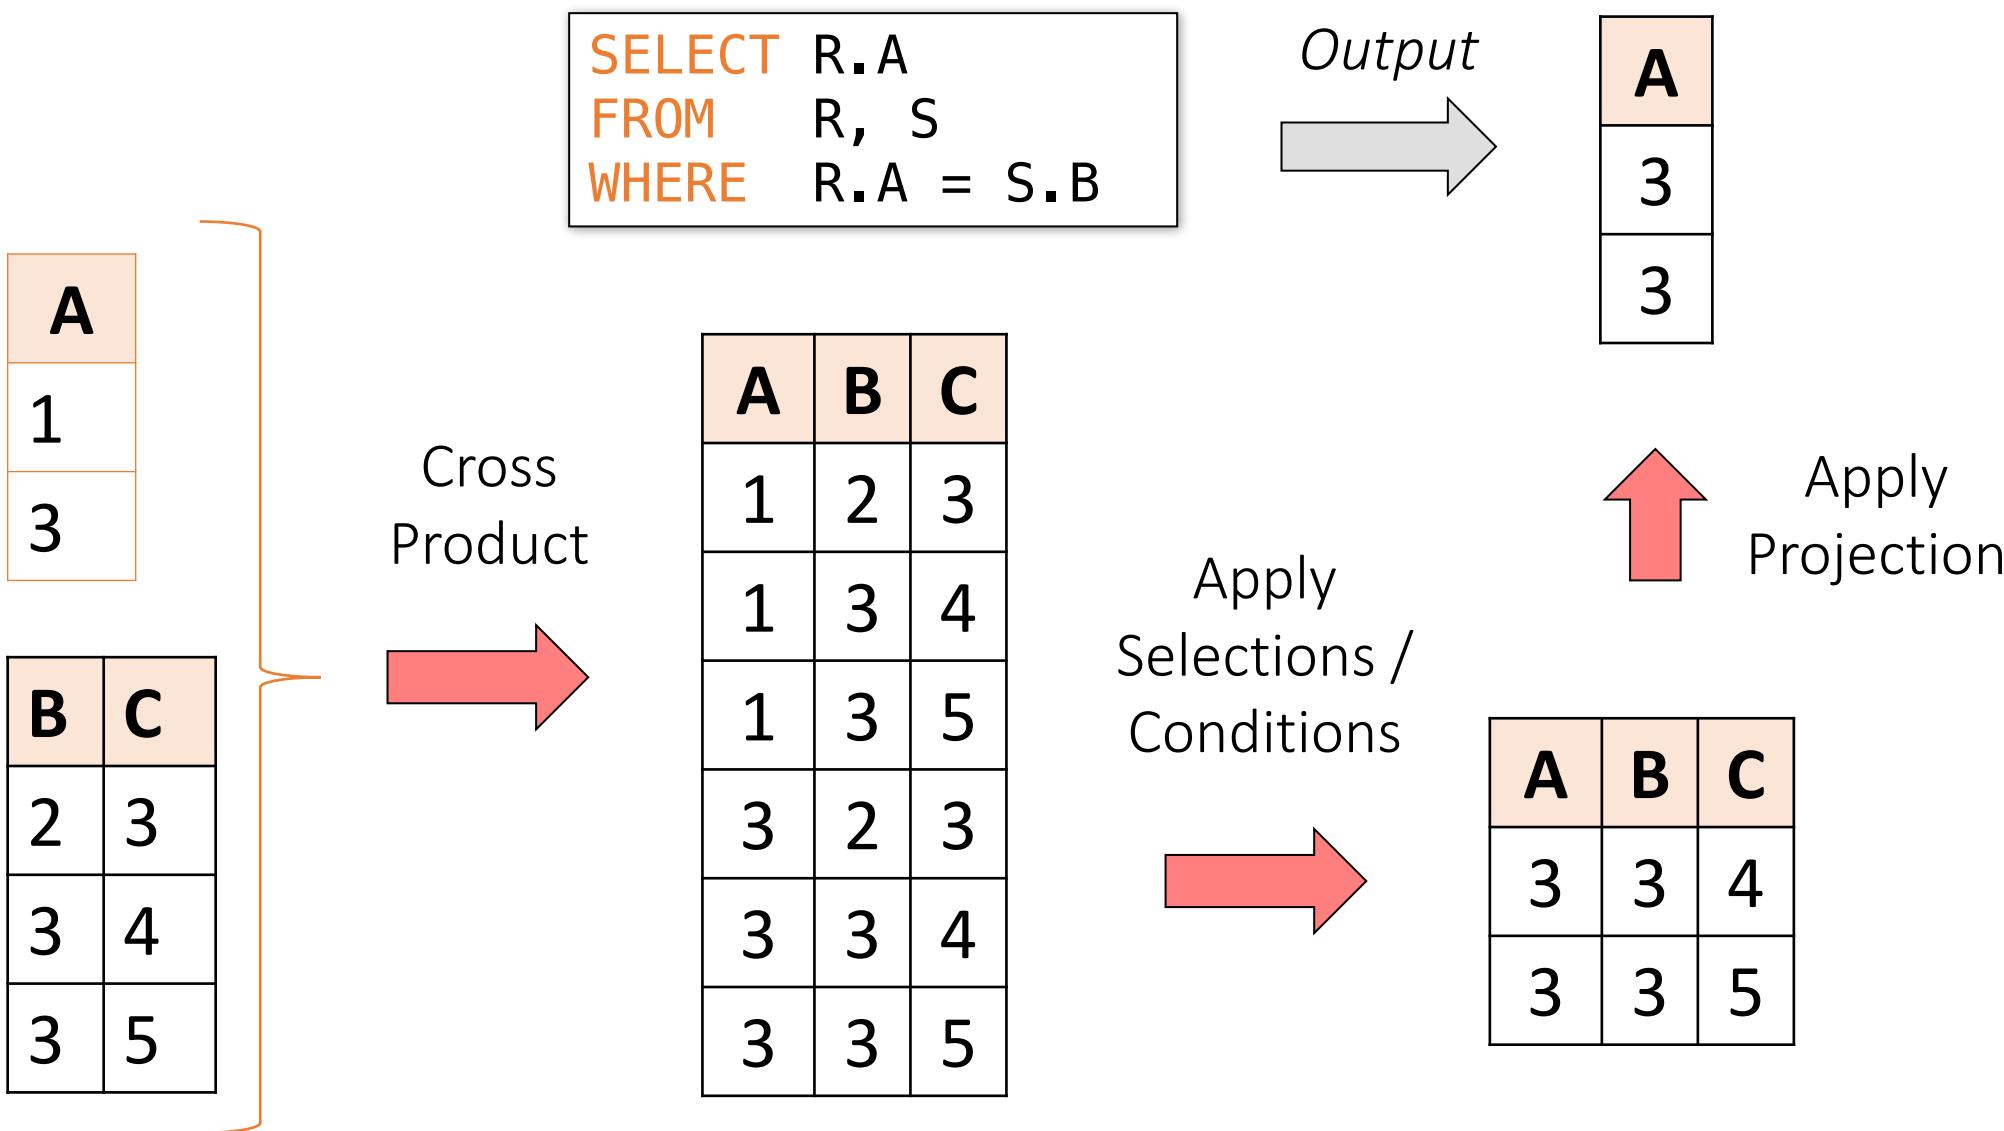


```

SELECT PName, Price
FROM Product, Company
WHERE Manufacturer = CName
AND Country='Japan'
AND Price <= 200
    
```

PName	Price
SingleTouch	\$149.99

An example of SQL semantics



Note the *semantics* of a join

```
SELECT R.A
FROM R, S
WHERE R.A = S.B
```

1. Take cross product:

$$X = R \times S$$

Recall: Cross product ($A \times B$) is the set of all unique tuples in A, B

Ex: $\{a,b,c\} \times \{1,2\}$
 $= \{(a,1), (a,2), (b,1), (b,2), (c,1), (c,2)\}$

2. Apply selections / conditions:

$$Y = \{(r,s) \in X \mid r.A == r.B\}$$

= Filtering!

3. Apply projections to get final output:

$$Z = (y.A,) \text{ for } y \in Y$$

= Returning only *some* attributes

Remembering this order is critical to understanding the output of certain queries (see later on...)

Note: we say “semantics” not “execution order”

- The preceding slides show *what a join means*
- Not actually how the DBMS executes it under the covers

A Subtlety about Joins

```
Product(PName, Price, Category, Manufacturer)  
Company(CName, StockPrice, Country)
```

Find all countries that manufacture some product
in the ‘Gadgets’ category.

```
SELECT Country  
FROM Product, Company  
WHERE Manufacturer=CName AND Category='Gadgets'
```

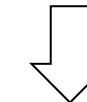
A subtlety about Joins

Product

PName	Price	Category	Manuf
Gizmo	\$19	Gadgets	GWorks
Powergizmo	\$29	Gadgets	GWorks
SingleTouch	\$149	Photography	Canon
MultiTouch	\$203	Household	Hitachi

Company

Cname	Stock	Country
GWorks	25	USA
Canon	65	Japan
Hitachi	15	Japan



```
SELECT Country
FROM Product, Company
WHERE Manufacturer=Cname
AND Category='Gadgets'
```

Country
?
?

What is the problem ?
What's the solution ?

ACTIVITY: [Lecture-2-3.ipynb](#)

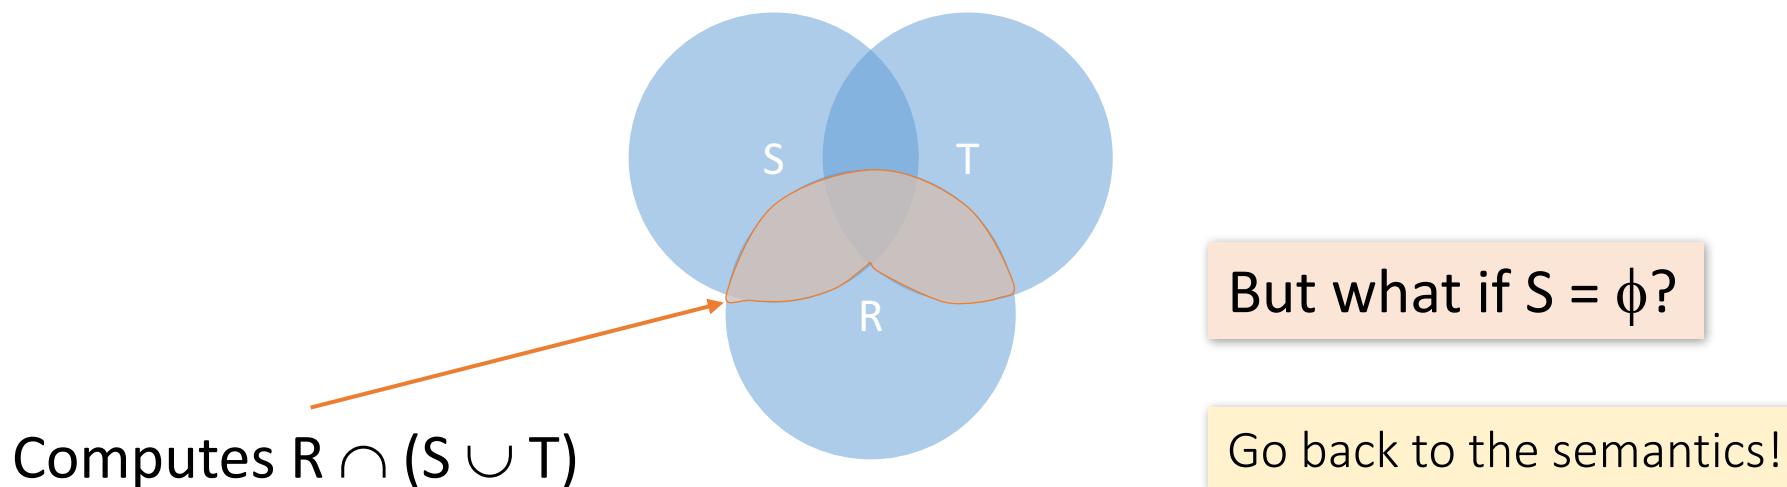
An Unintuitive Query

```
SELECT DISTINCT R.A  
FROM   R, S, T  
WHERE  R.A=S.A OR R.A=T.A
```

What does it compute?

An Unintuitive Query

```
SELECT DISTINCT R.A  
FROM   R, S, T  
WHERE  R.A=S.A OR R.A=T.A
```



An Unintuitive Query

```
SELECT DISTINCT R.A  
FROM   R, S, T  
WHERE  R.A=S.A OR R.A=T.A
```

- Recall the semantics!
 1. Take cross-product
 2. Apply selections / conditions
 3. Apply projection
- If $S = \{\}$, then the cross product of $R, S, T = \{\}$, and the query result = $\{\}$!

Must consider semantics here.

Are there more explicit way to do set operations like this?

Lecture 3: Advanced SQL – Part I

1. Set Operators & Nested Queries

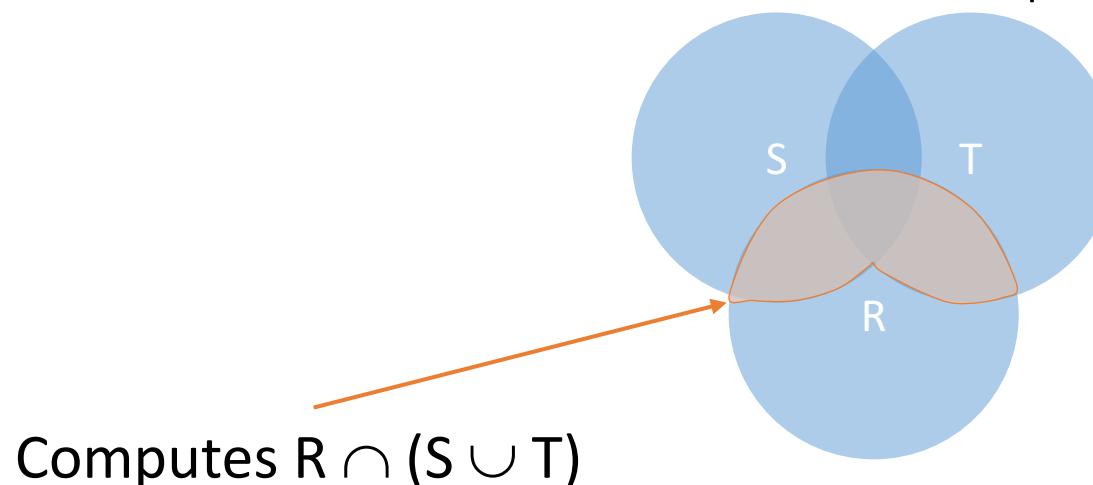
What you will learn about in this section

1. Multiset operators in SQL
2. Nested queries
3. ACTIVITY: Set operator subtleties

An Unintuitive Query

```
SELECT DISTINCT R.A  
FROM   R, S, T  
WHERE  R.A=S.A OR R.A=T.A
```

What does it compute?



But what if $S = \phi$?

Go back to the semantics!

An Unintuitive Query

```
SELECT DISTINCT R.A  
FROM   R, S, T  
WHERE  R.A=S.A OR R.A=T.A
```

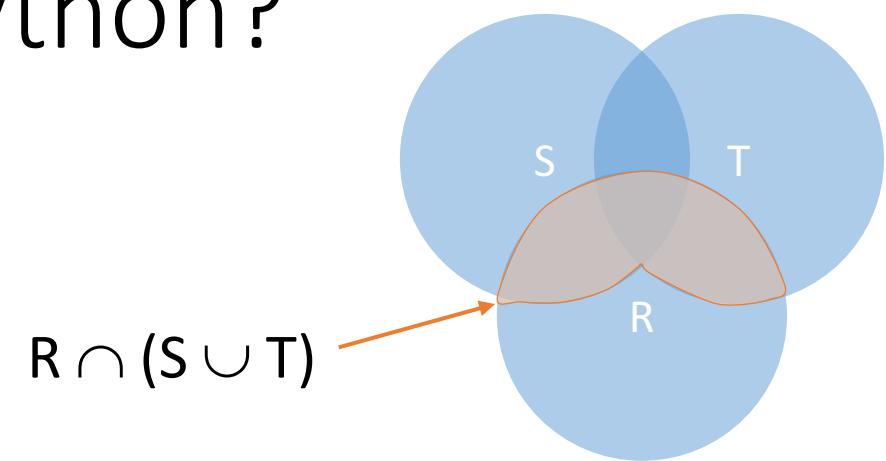
- Recall the semantics!
 1. Take cross-product
 2. Apply selections / conditions
 3. Apply projection
- If $S = \{\}$, then the cross product of $R, S, T = \{\}$, and the query result = $\{\}$!

Must consider semantics here.

Are there more explicit way to do set operations like this?

What does this look like in Python?

```
SELECT DISTINCT R.A  
FROM   R, S, T  
WHERE  R.A=S.A OR R.A=T.A
```



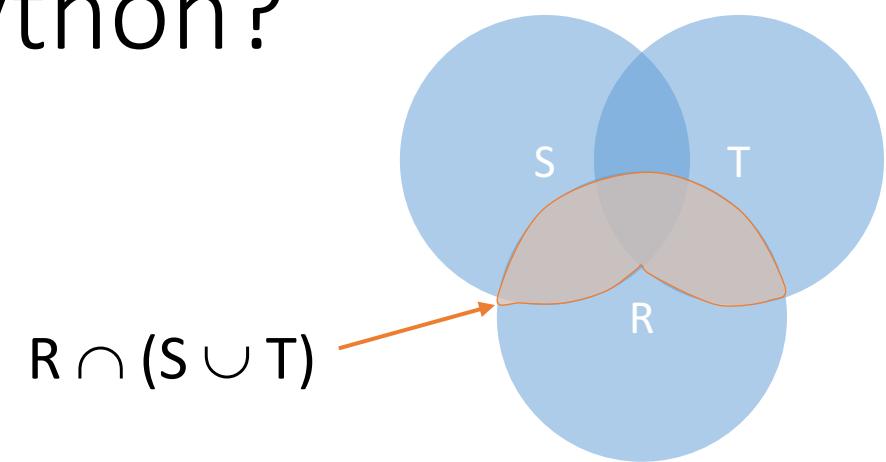
- Semantics:
 1. Take cross-product
 2. Apply selections / conditions
 3. Apply projection

Joins / cross-products are just nested for loops (in simplest implementation)!

If-then statements!

What does this look like in Python?

```
SELECT DISTINCT R.A  
FROM   R, S, T  
WHERE  R.A=S.A OR R.A=T.A
```



```
output = []  
  
for r in R:  
    for s in S:  
        for t in T:  
            if r['A'] == s['A'] or r['A'] == t['A']:  
                output.append(r['A'])  
return list(output)
```

Can you see now what happens if $S = []$?

Multiset Operations

Recall Multisets

Multiset X

Tuple
(1, a)
(1, a)
(1, b)
(2, c)
(2, c)
(2, c)
(1, d)
(1, d)



Equivalent
Representations
of a Multiset

$\lambda(X)$ = “Count of tuple in X”
(Items not listed have implicit count 0)

Multiset X

Tuple	$\lambda(X)$
(1, a)	2
(1, b)	1
(2, c)	3
(1, d)	2

Note: In a set all counts are {0,1}.

Generalizing Set Operations to Multiset Operations

Multiset X

Tuple	$\lambda(X)$
(1, a)	2
(1, b)	0
(2, c)	3
(1, d)	0

Multiset Y

Tuple	$\lambda(Y)$
(1, a)	5
(1, b)	1
(2, c)	2
(1, d)	2

 \cap $=$

Multiset Z

Tuple	$\lambda(Z)$
(1, a)	2
(1, b)	0
(2, c)	2
(1, d)	0

$$\lambda(Z) = \min(\lambda(X), \lambda(Y))$$

For sets, this is
intersection

Generalizing Set Operations to Multiset Operations

Multiset X

Tuple	$\lambda(X)$
(1, a)	2
(1, b)	0
(2, c)	3
(1, d)	0

Multiset Y

Tuple	$\lambda(Y)$
(1, a)	5
(1, b)	1
(2, c)	2
(1, d)	2

 \cup $=$

Multiset Z

Tuple	$\lambda(Z)$
(1, a)	5
(1, b)	1
(2, c)	3
(1, d)	2

$$\lambda(Z) = \max(\lambda(X), \lambda(Y))$$

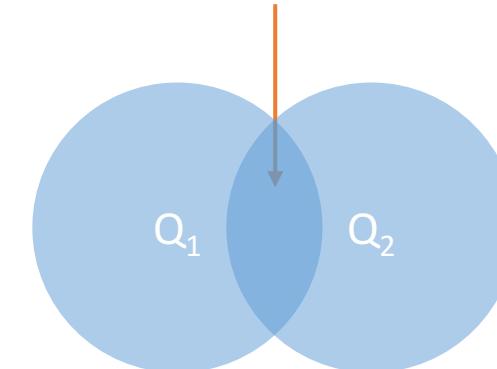
For sets,
this is **union**

Multiset Operations in SQL

Explicit Set Operators: INTERSECT

```
SELECT R.A  
FROM   R, S  
WHERE  R.A=S.A  
INTERSECT  
SELECT R.A  
FROM   R, T  
WHERE  R.A=T.A
```

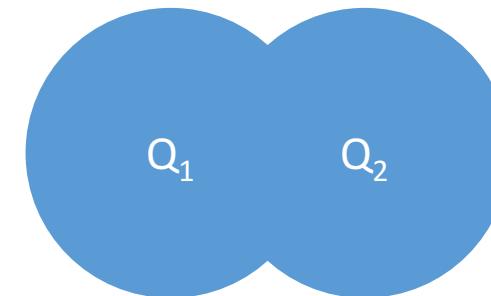
$$\{r.A \mid r.A = s.A\} \cap \{r.A \mid r.A = t.A\}$$



UNION

```
SELECT R.A  
FROM R, S  
WHERE R.A=S.A  
UNION  
SELECT R.A  
FROM R, T  
WHERE R.A=T.A
```

$$\{r.A \mid r.A = s.A\} \cup \{r.A \mid r.A = t.A\}$$



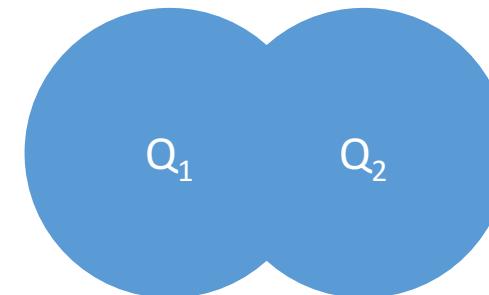
Why aren't there
duplicates?

What if we want
duplicates?

UNION ALL

```
SELECT R.A  
FROM R, S  
WHERE R.A=S.A  
UNION ALL  
SELECT R.A  
FROM R, T  
WHERE R.A=T.A
```

$$\{r.A \mid r.A = s.A\} \cup \{r.A \mid r.A = t.A\}$$

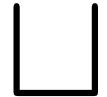


ALL indicates the Multiset disjoint union operation

Generalizing Set Operations to Multiset Operations

Multiset X

Tuple	$\lambda(X)$
(1, a)	2
(1, b)	0
(2, c)	3
(1, d)	0



Multiset Y

Tuple	$\lambda(Y)$
(1, a)	5
(1, b)	1
(2, c)	2
(1, d)	2



Multiset Z

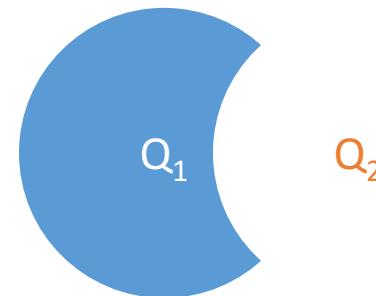
Tuple	$\lambda(Z)$
(1, a)	7
(1, b)	1
(2, c)	5
(1, d)	2

$$\lambda(Z) = \lambda(X) + \lambda(Y)$$

For sets,
this is **disjoint**
union

EXCEPT

```
SELECT R.A  
FROM   R, S  
WHERE  R.A=S.A  
EXCEPT  
SELECT R.A  
FROM   R, T  
WHERE  R.A=T.A
```

$$\{r.A \mid r.A = s.A\} \setminus \{r.A \mid r.A = t.A\}$$


What is the multiset version?

$$\lambda(Z) = \lambda(X) - \lambda(Y)$$

For elements that are in X

INTERSECT: Still some subtle problems...

```
Company(name, hq_city)  
Product(pname, maker, factory_loc)
```

```
SELECT hq_city  
FROM Company, Product  
WHERE maker = name  
      AND factory_loc = 'US'  
INTERSECT  
SELECT hq_city  
FROM Company, Product  
WHERE maker = name  
      AND factory_loc = 'China'
```

“Headquarters of companies which make gizmos in US AND China”

What if two companies have HQ in US: BUT one has factory in China (but not US) and vice versa? **What goes wrong?**

INTERSECT: Remember the semantics!

```
Company(name, hq_city) AS C
Product(pname, maker,
factory_loc) AS P
```

```
SELECT hq_city
FROM Company, Product
WHERE maker = name
AND factory_loc='US'
```

INTERSECT

```
SELECT hq_city
FROM Company, Product
WHERE maker = name
AND factory_loc='China'
```

Example: C JOIN P on maker = name

C.name	C.hq_city	P.pname	P.maker	P.factory_loc
X Co.	Seattle	X	X Co.	U.S.
Y Inc.	Seattle	X	Y Inc.	China

INTERSECT: Remember the semantics!

```
Company(name, hq_city) AS C
Product(pname, maker,
factory_loc) AS P
```

```
SELECT hq_city
FROM Company, Product
WHERE maker = name
AND factory_loc='US'
```

INTERSECT

```
SELECT hq_city
```

```
FROM Company, Product
WHERE maker = name
AND factory_loc='China'
```

Example: C JOIN P on maker = name

C.name	C.hq_city	P.pname	P.maker	P.factory_loc
X Co.	Seattle	X	X Co.	U.S.
Y Inc.	Seattle	X	Y Inc.	China

X Co has a factory in the US (but not China)
 Y Inc. has a factor in China (but not US)

But Seattle is returned by the query!

We did the INTERSECT
 on the wrong attributes!

One Solution: Nested Queries

```
Company(name, hq_city)  
Product(pname, maker, factory_loc)
```

```
SELECT DISTINCT hq_city  
FROM Company, Product  
WHERE maker = name  
AND name IN (  
    SELECT maker  
    FROM Product  
    WHERE factory_loc = 'US')  
AND name IN (  
    SELECT maker  
    FROM Product  
    WHERE factory_loc = 'China')
```

“Headquarters of companies which make gizmos in US AND China”

Note: If we hadn't used DISTINCT here, how many copies of each hq_city would have been returned?

High-level note on nested queries

- We can do nested queries because SQL is ***compositional***:
 - Everything (inputs / outputs) is represented as multisets- the output of one query can thus be used as the input to another (nesting)!
- This is extremely powerful!

Nested queries: Sub-queries Returning Relations

Another example:

```
Company(name, city)
Product(name, maker)
Purchase(id, product, buyer)
```

```
SELECT c.city
FROM Company c
WHERE c.name IN (
    SELECT pr.maker
    FROM Purchase p, Product pr
    WHERE p.product = pr.name
    AND p.buyer = 'Joe Blow')
```

“Cities where one can find companies that manufacture products bought by Joe Blow”

Nested Queries

Is this query equivalent?

```
SELECT c.city  
FROM Company c,  
      Product pr,  
      Purchase p  
WHERE c.name = pr.maker  
  AND pr.name = p.product  
  AND p.buyer = 'Joe Blow'
```

Beware of duplicates!

Nested Queries

```
SELECT DISTINCT c.city
FROM Company c,
      Product pr,
      Purchase p
WHERE c.name = pr.maker
  AND pr.name = p.product
  AND p.buyer = 'Joe Blow'
```

```
SELECT DISTINCT c.city
FROM Company c
WHERE c.name IN (
    SELECT pr.maker
    FROM Purchase p, Product pr
    WHERE p.product = pr.name
        AND p.buyer = 'Joe Blow')
```

Now they are equivalent

Subqueries Returning Relations

You can also use operations of the form:

- s > ALL R
- s < ANY R
- EXISTS R

ANY and ALL not supported by
SQLite.

Ex: **Product(name, price, category, maker)**

```
SELECT name
FROM Product
WHERE price > ALL(
    SELECT price
    FROM Product
    WHERE maker = 'Gizmo-Works')
```

Find products that
are more expensive
than all those
produced by
“Gizmo-Works”

Subqueries Returning Relations

You can also use operations of the form:

- $s > \text{ALL } R$
- $s < \text{ANY } R$
- EXISTS R

Ex: **Product(name, price, category, maker)**

```
SELECT p1.name
FROM Product p1
WHERE p1.maker = 'Gizmo-Works'
AND EXISTS(
    SELECT p2.name
    FROM Product p2
    WHERE p2.maker <> 'Gizmo-Works'
    AND p1.name = p2.name)
```

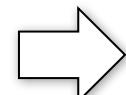
<> means \neq

Find ‘copycat’ products, i.e. products made by competitors with the same names as products made by “Gizmo-Works”

Nested queries as alternatives to INTERSECT and EXCEPT

INTERSECT and EXCEPT not in some DBMSs!

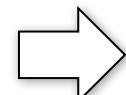
```
(SELECT R.A, R.B
FROM R)
INTERSECT
(SELECT S.A, S.B
FROM S)
```



```
SELECT R.A, R.B
FROM R
WHERE EXISTS(
    SELECT *
    FROM S
    WHERE R.A=S.A AND R.B=S.B)
```

If R, S have no duplicates, then can write without sub-queries (HOW?)

```
(SELECT R.A, R.B
FROM R)
EXCEPT
(SELECT S.A, S.B
FROM S)
```



```
SELECT R.A, R.B
FROM R
WHERE NOT EXISTS(
    SELECT *
    FROM S
    WHERE R.A=S.A AND R.B=S.B)
```

Correlated Queries

```
Movie(title, year, director, length)
```

```
SELECT DISTINCT title
FROM Movie AS m
WHERE year <> ANY(
    SELECT year
    FROM Movie
    WHERE title = m.title)
```

Find movies whose title appears more than once.

Note the scoping of the variables!

Note also: this can still be expressed as single SFW query...

Complex Correlated Query

Product(name, price, category, maker, year)

```
SELECT DISTINCT x.name, x.maker
FROM Product AS x
WHERE x.price > ALL(
    SELECT y.price
    FROM Product AS y
    WHERE x.maker = y.maker
    AND y.year < 1972)
```

Find products (and their manufacturers) that are more expensive than all products made by the same manufacturer before 1972

Can be very powerful (also much harder to optimize)

[Activity-3-1.ipynb](#)

Basic SQL Summary

- SQL provides a high-level declarative language for manipulating data (DML)
- The workhorse is the SFW block
- Set operators are powerful but have some subtleties
- Powerful, nested queries also allowed.