

Lecture 20: Transactions

Announcements

- Most important lesson from Project 3!??!!
- Project 2 grades should be up
 - Let us know of questions
- Project 4 coming out later today:
 - Re-use your database from Project 1 and add a nice web interface on top to support transactions.
- Evaluation should be open please do not forget
- December 6th: Guest lecture by Stratis Viglas from Google
 - ATTEND!!!



Goals for this pair of lectures

- **Transactions** are a programming abstraction that enables the DBMS to handle *recovery* and *concurrency* for users.
- **Application:** Transactions are critical for users
 - Even casual users of data processing systems!
- **Fundamentals:** The basics of **how TXNs work**
 - Transaction processing is part of the debate around new data processing systems
 - Give you enough information to understand how TXNs work, and the main concerns with using them

Lecture 20: Intro to Transactions & Logging

Today's Lecture

1. Transactions
2. Properties of Transactions: ACID
3. Logging

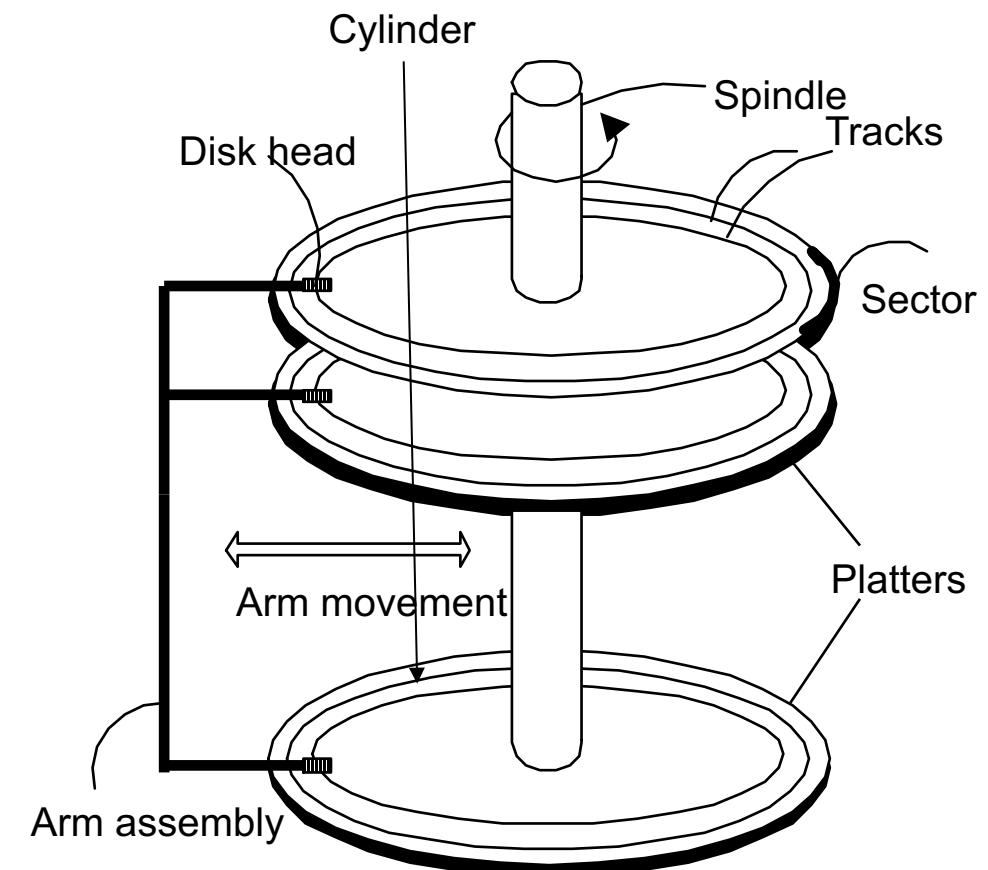
1. Transactions

What you will learn about in this section

1. Our “model” of the DBMS / computer
2. Transactions basics
3. Motivation: Recovery & Durability
4. Motivation: Concurrency [*next lecture*]

High-level: Disk vs. Main Memory

- **Disk:**
 - *Slow*
 - Sequential access
 - (although fast sequential reads)
 - *Durable*
 - We will assume that once on disk, data is safe!
 - *Cheap*



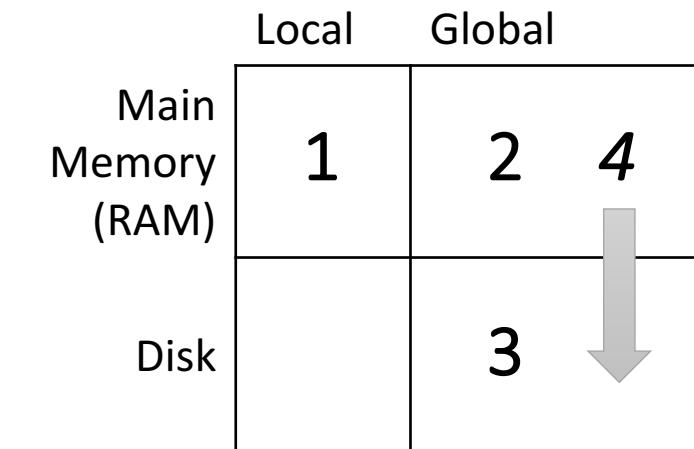
High-level: Disk vs. Main Memory

- Random Access Memory (RAM) or **Main Memory**:
 - *Fast*
 - Random access, byte addressable
 - ~10x faster for sequential access
 - ~100,000x faster for random access!
 - *Volatile*
 - Data can be lost if e.g. crash occurs, power goes out, etc!
 - *Expensive*
 - For \$100, get 16GB of RAM vs. 2TB of disk!



Our model: Three Types of Regions of Memory

1. **Local:** In our model each process in a DBMS has its own local memory, where it stores values that only it “sees”
2. **Global:** Each process can read from / write to shared data in main memory
3. **Disk:** Global memory can read from / flush to disk
4. **Log:** Assume on stable disk storage- spans both main memory and disk...



Log is a *sequence* from main memory -> disk

“Flushing to disk” = writing to disk from main memory

High-level: Disk vs. Main Memory

- Keep in mind the tradeoffs here as motivation for the mechanisms we introduce
 - Main memory: fast but limited capacity, volatile
 - Vs. Disk: slow but large capacity, durable

How do we effectively utilize *both* ensuring certain critical guarantees?

Transactions

Transactions: Basic Definition

A transaction (“TXN”) is a sequence of one or more *operations* (reads or writes) which reflects *a single real-world transition*.

In the real world, a TXN either happened completely or not at all

```
START TRANSACTION  
    UPDATE Product  
        SET Price = Price - 1.99  
        WHERE pname = 'Gizmo'  
    COMMIT
```

Transactions: Basic Definition

A transaction (“TXN”) is a sequence of one or more *operations* (reads or writes) which reflects a *single real-world transition*.

In the real world, a TXN either happened completely or not at all

Examples:

- Transfer money between accounts
- Purchase a group of products
- Register for a class (either waitlist or allocated)

Transactions in SQL

- In “ad-hoc” SQL:
 - Default: each statement = one transaction
- In a program, multiple statements can be grouped together as a transaction:

```
START TRANSACTION
    UPDATE Bank SET amount = amount - 100
    WHERE name = 'Bob'
    UPDATE Bank SET amount = amount + 100
    WHERE name = 'Joe'
COMMIT
```

Model of Transaction for this Class

Note: We assume that the DBMS *only* sees reads and writes to data

- User may do much more
- In real systems, databases do have more info...

Motivation for Transactions

Grouping user actions (reads & writes) into *transactions* helps with two goals:

1. **Recovery & Durability**: Keeping the DBMS data consistent and durable in the face of crashes, aborts, system shutdowns, etc.
2. **Concurrency**: Achieving better performance by parallelizing TXNs *without* creating anomalies

This lecture!

Next lecture

Motivation

1. Recovery & Durability of user data is essential for reliable DBMS usage

- The DBMS may experience crashes (e.g. power outages, etc.)
- Individual TXNs may be aborted (e.g. by the user)

Idea: Make sure that TXNs are either durably stored in full, or not at all; keep log to be able to “roll-back” TXNs

Protection against crashes / aborts

Client 1:

```
INSERT INTO SmallProduct(name, price)  
    SELECT pname, price  
        FROM Product  
    WHERE price <= 0.99
```

Crash / abort!

```
DELETE Product  
    WHERE price <=0.99
```

What goes wrong?

Protection against crashes / aborts

Client 1:

```
START TRANSACTION
    INSERT INTO SmallProduct(name, price)
        SELECT pname, price
        FROM Product
        WHERE price <= 0.99

    DELETE Product
        WHERE price <=0.99
    COMMIT OR ROLLBACK
```

Now we'd be fine! We'll see how / why this lecture

Motivation

2. Concurrent execution of user programs is essential for good DBMS performance.

- Disk accesses may be frequent and **slow**- optimize for throughput (# of TXNs), trade for latency (time for any one TXN)
- Users should still be able to execute TXNs as if in **isolation** and such that **consistency** is maintained

Idea: Have the DBMS handle running several user TXNs concurrently, in order to keep CPUs humming...

Multiple users: single statements

```
Client 1: UPDATE Product  
          SET Price = Price - 1.99  
          WHERE pname = 'Gizmo'
```

```
Client 2: UPDATE Product  
          SET Price = Price*0.5  
          WHERE pname='Gizmo'
```

Two managers attempt to discount products *concurrently*-
What could go wrong?

Multiple users: single statements

```
Client 1: START TRANSACTION
          UPDATE Product
          SET Price = Price - 1.99
          WHERE pname = 'Gizmo'
          COMMIT

Client 2: START TRANSACTION
          UPDATE Product
          SET Price = Price*0.5
          WHERE pname='Gizmo'
          COMMIT
```

Now works like a charm- we'll see how / why next lecture...

2. Properties of Transactions

What you will learn about in this section

1. Atomicity
2. Consistency
3. Isolation
4. Durability
5. ACTIVITY?

Transaction Properties: ACID

- **Atomic**
 - State shows either all the effects of txn, or none of them
- **Consistent**
 - Txn moves from a state where integrity holds, to another where integrity holds
- **Isolated**
 - Effect of txns is the same as txns running one after another (ie looks like batch mode)
- **Durable**
 - Once a txn has committed, its effects remain in the database

ACID continues to be a source of great debate!

ACID: Atomicity

- TXN's activities are atomic: **all or nothing**
 - Intuitively: in the real world, a transaction is something that would either occur *completely* or *not at all*
- Two possible outcomes for a TXN
 - It *commits*: all the changes are made
 - It *aborts*: no changes are made

ACID: Consistency

- The tables must always satisfy user-specified ***integrity constraints***
 - *Examples:*
 - Account number is unique
 - Stock amount can't be negative
 - Sum of *debits* and of *credits* is 0
- How consistency is achieved:
 - Programmer makes sure a txn takes a consistent state to a consistent state
 - *System* makes sure that the txn is **atomic**

ACID: Isolation

- A transaction executes concurrently with other transactions
- **Isolation:** the effect is as if each transaction executes in *isolation* of the others.
 - E.g. Should not be able to observe changes from other transactions during the run

ACID: Durability

- The effect of a TXN must continue to exist (“*persist*”) after the TXN
 - And after the whole program has terminated
 - And even if there are power failures, crashes, etc.
 - And etc...
- Means: Write data to **disk**

Change on the horizon?
Non-Volatile Ram (NVRam).
Byte addressable.

Challenges for ACID properties

- In spite of failures: Power failures, but not media failures
- Users may abort the program: need to “rollback the changes”
 - Need to *log* what happened
- Many users executing concurrently
 - Can be solved via locking (we’ll see this next lecture!)

This lecture

Next lecture

And all this with... Performance!!

A Note: ACID is contentious!



- Many debates over ACID, both **historically** and **currently**
- Many newer “NoSQL” DBMSs relax ACID
- In turn, now “NewSQL” reintroduces ACID compliance to NoSQL-style DBMSs...



ACID is an extremely important & successful paradigm, but still debated!

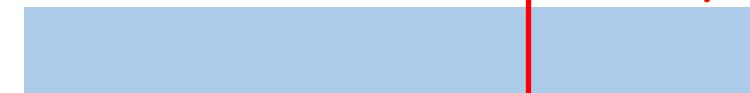
Goal for this lecture: Ensuring Atomicity & Durability

ACID

- Atomicity:

- TXNs should either happen completely or not at all
- If abort / crash during TXN, *no* effects should be seen

TXN 1



No changes persisted

- Durability:

- If DBMS stops running, changes due to completed TXNs should all persist
- *Just store on stable disk*

TXN 2



All changes persisted

We'll focus on how to accomplish atomicity (via logging)

The Log

- Is a list of modifications
- Log is *duplexed* and *archived* on stable storage.
- Can **force write** entries to disk
 - A page goes to disk.
- All log activities *handled transparently* the DBMS.

Assume we
don't lose it!

Basic Idea: (Physical) Logging

- Record UNDO information for every update!
 - Sequential writes to log
 - Minimal info (diff) written to log
- The **log** consists of an **ordered list of actions**
 - Log record contains:
<XID, location, old data, new data>

This is sufficient to UNDO any transaction!

Why do we need logging for atomicity?

- Couldn't we just write TXN to disk **only** once whole TXN complete?
 - Then, if abort / crash and TXN not complete, it has no effect- atomicity!
 - *With unlimited memory and time, this could work...*
- However, we **need to log partial results of TXNs** because of:
 - Memory constraints (enough space for full TXN??)
 - Time constraints (what if one TXN takes very long?)

We need to write partial results to disk!

...And so we need a **log** to be able to *undo* these partial results!

3. Atomicity & Durability via Logging

What you will learn about in this section

1. Logging: An animation of commit protocols

A Picture of Logging

A picture of logging

$T: R(A), W(A)$



A picture of logging

$T: R(A), W(A)$

$A: 0 \rightarrow 1$



A picture of logging

$T: R(A), W(A)$

$A: 0 \rightarrow 1$



Operation recorded in log in main memory!



What is the correct way to write this all to disk?

- We'll look at the *Write-Ahead Logging (WAL)* protocol
- We'll see why it works by looking at other protocols which are incorrect!

Remember: Key idea is to ensure durability
while maintaining our ability to “undo”!

Write-Ahead Logging (WAL) TXN Commit Protocol

Transaction Commit Process

1. FORCE Write **commit** record to log
2. All log records up to last update from this TX are FORCED
3. Commit() returns

Transaction is committed *once commit log record is on stable storage*

Incorrect Commit Protocol #1

T: R(A), W(A)

A: 0 → 1



Let's try committing
before we've written
either data or log to
disk...

OK, Commit!

If we crash now, is T
durable?

Lost T's update!

Incorrect Commit Protocol #2

T: R(A), W(A)

A: 0 → 1



Let's try committing
after we've written
data but *before* we've
written log to disk...

OK, Commit!

If we crash now, is T
durable? Yes! Except...

**How do we know
whether T was
committed??**

Improved Commit Protocol (WAL)

Write-ahead Logging (WAL) Commit Protocol

T: R(A), W(A)



This time, let's try committing after we've written log to disk but before we've written data to disk... this is WAL!

OK, Commit!

If we crash now, is T durable?

Write-ahead Logging (WAL) Commit Protocol

T: R(A), W(A)

T →



A: 0 → 1



This time, let's try committing after we've written log to disk but before we've written data to disk... this is WAL!

OK, Commit!

If we crash now, is T durable?

USE THE LOG!

Write-Ahead Logging (WAL)

- DB uses **Write-Ahead Logging (WAL)** Protocol:

Each update is logged! Why not reads?

1. Must *force log record* for an update *before* the corresponding data page goes to storage

→ Atomicity

2. Must *write all log records* for a TX *before commit*

→ Durability

Logging Summary

- If DB says TX **commits**, TX effect **remains** after database crash
- DB can **undo actions** and help us with **atomicity**
- This is only half the story...